

## COSC3081/3015 Advanced Programming for Data Science

### Assignment 3: NLP Web-based Data Application

#### Milestone I: Natural Language Processing

Assessment Type	Group assignment. Submit online via Canvas→ Assignments→ Assignment 3→ Milestone I: Natural Language Processing. Marks are awarded for meeting requirements as closely as possible. Clarifications/updates may be made via announcements/relevant discussion forums.
Due Date	End of Week 10, Please check Canvas for the exact due date and time
Marks	20

### 1. Overview

Today, many online clothing websites provide the convenience of purchasing clothes remotely, eliminating the need to visit physical stores. These sites often incorporate recommendation systems that suggest similar products to customers, aiming to boost sales by offering items tailored to their preferences. Product recommendations and sales data analysis are becoming increasingly important and beneficial for e-commerce platforms, as they can significantly enhance the shopping experience for customers and improve the overall efficiency of online shopping systems.

With advances in text analysis, the reviews by customers sharing similar shopping histories can help in developing a classification model for recommending products to similar customers. This is a strategy commonly used by online shopping websites to increase the sales of the products. To do so, we need an automated classification system that helps to predict whether the customer will recommend the product or not.

This assessment includes two milestones. The first milestone (NLP) concerns the pipeline from basic text preprocessing to building text classification models for predicting the recommendation of a clothing item. Then, the second milestone will adopt one of the models that we built in the first milestone and develop an online shopping website that allows users to browse existing products, as well as identify the recommendations of a product.

This assessment description is about Milestone 1: Natural Language Processing.

**Data source:** <https://www.kaggle.com/datasets/nicapotato/womens-ecommerce-clothing-reviews>

**Note:** The dataset has been modified for the course, hence, it is not in its original state

### 2. Learning Outcomes

This assessment relates to the following learning outcomes of the course:

- CLO 4: Pre-process natural language text data to generate effective feature representations;
- CLO 5: Document and maintain an editable transcript of the data pre-processing pipeline for professional reporting.

### 3. Assessment Details

In this milestone, you are required to pre-process a collection of clothing reviews, build machine learning models for recommendation (i.e., classifying whether the review represents the recommendation of a clothing item or not), and perform evaluation and analysis on the built models.

## The Data

In this assignment, you are given a collection of clothing reviews (approx 19600 reviews). The data folder is available to download from Canvas. The three features that you will use in this assignment are: “**Title**” (title of the review), “**Review Text**” (the detailed review of the clothing item), and “**Recommended**” (label represent that the clothing item is recommended by the reviewer or not. ‘0’ and ‘1’ represent “*not-recommended*” and “*recommended*” respectively).

## Task 1: Basic Text Pre-processing [5 marks]

In this task, you are required to perform basic text pre-processing on the given dataset, including, but not limited to tokenization, removing most/least frequent words and stop words. **In this task, we focus on pre-processing the “Review Text” only.** You are required to perform the following:

1. Extract information about the review. Perform the pre-processing steps mentioned below to the extracted reviews
2. Tokenize each clothing review. The word tokenization must use the following regular expression, **r"[a-zA-Z]+(?:['-']+[a-zA-Z]+)?"**;
3. All the words must be converted into the lower case;
4. Remove words with a length less than 2.
5. Remove stopwords using the provided stop words list (i.e., **stopwords\_en.txt**). It is located inside the same downloaded folder.
6. Remove the word that appears only once in the document collection, based on **term** frequency.
7. Remove the top 20 most frequent words based on **document** frequency.
8. Save the processed data as **processed.csv** file.
9. Build a vocabulary of the cleaned/processed reviews, and save it in a txt file (**please refer to the Required Output section**);

Note:

- For all the words that we removed (including step 4,5,6,7), you will also exclude them in the generated vocabulary.
- The output of this task will be checked against the expected output. You should strictly follow the following format requirement.

## Required Output for Task 1:

The output of this task must contain the following files:

- **vocab.txt** This file contains the **unigram** vocabulary, one each line, in the following format: **word\_string:word\_integer\_index**. Very importantly, **words in the vocabulary must be sorted in alphabetical order, and the index value starts from 0**. This file is the key to interpret the sparse encoding. For instance, in the following example, the word *aaron* is the 20th word (the corresponding integer\_index as 19) in the vocabulary (note that the index values and words in the following image are artificial and used to demonstrate the required format only, it doesn't reflect the values of the actual expected output).

```
aand:12
antrekkelijk:13
aap:14
aar:15
aarca:16
aardman:17
aareon:18
aaron:19
aaronwallis:20
aarolette:21
aarolettehotel:22
```

Fig.1 example format for vocab.txt

- All Python code related to Task 1 should be written in the jupyter notebook **task1.ipynb**.

## Task 2: Generating Feature Representations for Clothing Reviews [7 marks]

In this task, you are required to generate different types of feature representations for the clothing item reviews. **Note that in this task, we will only consider the **description** of the review (ignore *Title*).** The feature representation that you need to generate includes the following:

### Bag-of-words model:

- Generate the **Count** vector representation for each clothing review, and save them into a file (**please refer to the required output**). Note, **the generated Count vector representation must be based on the generated vocabulary in Task 1 (as saved in vocab.txt)**.

### Models based on word embeddings:

- You are required to generate a feature representation of clothing reviews based on the following language models:
  - Choose 1 embedding language model, e.g., FastText, GoogleNews300, or other Word2Vec pretrained models, or Glove, which will be used to extract the feature vector for each review.
  - You are required to build the **weighted (i.e., TF-IDF weighted) and unweighted** vector representation for each clothing review using the chosen language model.

To summarise, there are 3 different types of feature representation of documents that you need to build in this task, weighted, unweighted and embedding representation.

### Required Output for Task 2:

- **count\_vectors.txt** This file stores the sparse count vector representation of clothing reviews in the following format. Each line of this file corresponds to one review. It starts with a '#' key followed by the word index of the item review, and a comma ','. The rest of the line is the sparse representation of the corresponding description in the form of **word\_integer\_index:word\_freq** separated by comma. Following is an example of the file format (note that the following image is artificial and used to demonstrate the required format only, it doesn't reflect the values of the actual expected output):

```
#12612628,11:1,35:1,255:4,256:1,411:3,503:1,535:2,615:4,727:2,742:1,752:1,782
:1,833:3,862:1,893:1,941:1,1067:1
#12612830,0:1,64:1,70:1,72:1,90:1,92:1,116:1,218:1,231:1,245:1,270:1,274:1,31
0:1,316:1,317:1,330:2,412:2,433:1,441:1,529:1,532:1,537:1,538:1,545:2,616:2,6
17:1,649:1,680:1,697:1,737:1,742:1,744:1,752:1,766:1,768:1,773:2,774:1,799:1,
845:1,855:1,858:1,864:1,874:1,888:2,902:1,937:1,952:2,962:1,1003:1,1008:1,107
7:1,1082:1,1087:1,1092:1,1093:1
#12612844,6:2,17:1,57:1,59:2,134:2,149:2,169:2,196:1,198:1,199:2,208:3,214:1,
218:1,245:1,252:1,255:6,259:1,265:1,327:2,340:1,352:1,387:2,390:2,432:1,450:1
```

Fig. 4 Example format for count\_vectors.txt

Note: **word\_freq** here refers to the frequency of the unigram in the corresponding **Review Text** only, **excluding the Review Title**.

## Task 3: Clothing Review Classification [8 marks]

In this task, you are required to build machine learning models for classifying the recommendation of an item review text. A simple model that you can consider is the logistic regression model from sklearn as demonstrated in the activities. However, you feel free to select other models (even if it has not been covered in this course). You are required to conduct two sets of experiments on the provided dataset to investigate the following two questions, respectively.

### Q1: Language model comparisons

Which language model we built previously (based on clothing reviews) performs the best with the chosen machine learning model? To answer these questions, you are required to build machine learning models based on the feature representations of the documents you generated in **Task 2**, and to perform evaluation on the various model performance.

### Q2: Does more information provide higher accuracy?

In **Task 2**, we have built a number of feature representations of documents based on clothing reviews. However, we have not explored other features of an item review, e.g., the title of the review. Will adding extra information help to boost up the accuracy of the model? To answer this question, you are required to conduct experiments to build and **compare** the performance of classification models that consider:

- only title of the review
- only description/text of the review (which you've already done in Task 3)
- both title and detailed review (Review Text) of a clothing item. For this, you have the flexibility to simply concatenate title and description of an item review when generating document feature representation; or to generate separate feature representations for title and description, respectively, and use both features in the classification models.

**Note that:** For both questions above,

- Consider demonstrating the comparison using at least three different models: one based on the bag-of-words model, and the other two focusing on weighted and unweighted models.
- When evaluating the performance of the models, you are required to conduct a **5-fold cross validation** to obtain robust comparisons.

**All Python code related to Task 2 and 3 should be written in the jupyter notebook task2\_3.ipynb.**

## Marking Guidelines

### Marking Criteria

- **Mechanical pass:** Your outputs will be compared against the expected output. Therefore, marking will be based on the similarity between what we expect (as discussed in the instructions) and what we receive from you. It is extremely important to carefully follow the instructions to produce the expected output. Otherwise, you may easily lose many points for simple mistakes (e.g. typos in the format of the files, not loading essential libraries, different file names/path, etc).
- **Expert pass:** Your Jupyter notebook will be checked by an expert to validate the logic and flow, proper use of libraries and functions, and clarity of codes, comments, structure and presentation.
- You need to ensure all the codes and files that are required to run your code are included in the submission. The expert will **NOT** fix your code's problem even if it is a simple typo in a file name or an imported library.

### Mark Allocations

- Task 1 Basic Text Pre-processing [5%]
  - Implementation [4%]
  - Notebook presentation [1%], proportional to actual mark obtained in implementation
- Task 2 [7%]
  - Implementation [5%]
  - Notebook presentation [2%], proportional to actual mark obtained in implementation
- Task 3 [8%]

- Implementation [6%]
- Notebook presentation [2%], proportional to actual mark obtained in implementation

**For Task 1, and Task 2 and 3, you are required to maintain an auditable and editable transcript, and communicate any justification of methods/approach chosen, results, analysis and findings through jupyter notebook. The presentation of the jupyter notebook accounts for certain percentages of the allocated mark for each task, proportional to the actual mark obtained, as per specified above. Students can refer to the activities in modules as examples for the level of details that they should include in their jupyter notebook.**

## Submission

The final submission of this milestone will consist of:

- The required output from Task 1, including **vocab.txt**
- The required output from Task 2, **count\_vectors.txt**
- The jupyter notebook of Task 1, and Task 2&3, respectively
- The **.py** format of the jupyter notebook of Task 1, and Task 2&3, respectively. Note that:
  - the content of the .py file must match your jupyter notebook
  - the .py file will be used for plagiarism detections on both comment/description content, as well as the actual code.
  - to help promote academic integrity, please make sure you submit the .py files. Submission without the .py files or unmatched .py files will NOT be marked.
  - note that the .py files can be easily downloaded from jupyter notebook interface (File -> Download as -> Python (.py))
- Put all the above mentioned files in a folder, named with your student id, Zip the folder with the same name (i.e., **s1234567.zip**) and upload for submission

### Assessment declaration:

When you submit work electronically, you agree to the assessment declaration:

<https://www.rmit.edu.au/students/student-essentials/assessment-and-exams/assessment/assessment-declaration>

### Late Submission Penalty

Late submissions will incur a 10% penalty on the total marks of the corresponding assessment task per day or part of day late. Submissions that are late by 5 days or more are not accepted and will be awarded zero, unless special consideration has been granted. Granted Special Considerations with a new due date set more than 2 weeks after the original due will automatically result in an equivalent assessment in the form of a practical test with interview, assessing the same knowledge and skills of the assignment (location and time to be arranged by the instructor). Please ensure your submission is correct (all files are there, compiles etc), re-submissions after the due date and time will be considered as late submissions.

## Academic integrity and plagiarism (standard warning)

Academic integrity is about honest presentation of your academic work. It means acknowledging the work of others while developing your own insights, knowledge and ideas. You should take extreme care that you have:

- acknowledged words, data, diagrams, models, frameworks and/or ideas of others you have quoted (i.e. directly copied), summarised, paraphrased, discussed or mentioned in your assessment through the appropriate referencing methods,
- provided a reference list of the publication details so your reader can locate the source if necessary. This includes material taken from Internet sites.

If you do not acknowledge the sources of your material, you may be accused of plagiarism

because you have passed off the work and ideas of another person without appropriate referencing, as if they were your own.

RMIT University treats plagiarism as a very serious offence constituting misconduct. Plagiarism covers a variety of inappropriate behaviours, including:

- Failure to properly document a source
- Copyright material from the internet or databases
- Collusion between students

For further information on our policies and procedures, please refer to <https://www.rmit.edu.au/students/student-essentials/rights-and-responsibilities/academic-integrity>